Advertising in Toxic Feeds: Ad Effectiveness & Brand Safety

Registered Report Stage 1 Protocol

Jonas Goergen[2] & Hauke Roggenkamp[1]

[1]ETH Zurich, Department of Management, Technology, and Economics
[2]Kedge Business School, Department of Marketing

Toxic content on social media captures attention and drives engagement, yet brands overwhelmingly treat it as a reputational threat and withdraw advertising from toxic environments. These two responses reflect two literatures that examine the effects of toxicity in isolation: brand safety research documents associative contamination from unsafe adjacent content, while platform engagement research documents attentional and behavioral gains from negative and toxic content. No existing work examines both within the same feed exposure. We integrate these perspectives in a single empirical framework. Using a dose-response design with stimulus sampling, we manipulate the proportion of toxic content in simulated social media feeds and measure feed-level engagement, ad-level attention, and brand-level reputational outcomes. We predict an inverted-U relationship between feed-level toxicity and both feed engagement and ad engagement: moderate toxicity elevates arousal and deepens processing; high toxicity triggers aversion and disengagement. The reputational effect of toxicity on brand attitude is theoretically ambiguous — the same attentional gains that benefit the ad also increase exposure to the aversive context — and is examined as a research question. Our design independently varies feed-level toxicity dose and adjacent-post toxicity, allowing simultaneous estimation of global and local effects. The findings inform brand safety strategy by providing a basis for evaluating the net effect of feed-level toxicity on advertising effectiveness.

Keywords: toxicity, brand safety, advertising, platforms
Words: 7933

Introduction

Social media platforms have increasingly become home to toxic discourses. Harassment, hate speech, cyberbullying, identity attacks, and deliberately insulting language pervade feeds across topics and languages, driven by recommendation algorithms and virality mechanisms amplifying such content. This toxicity poses a challenge for the multi-billion dollar advertising industry operating on these platforms. When toxicity on X (formerly Twitter) escalated following changes in ownership and content moderation in late 2022, major advertisers including Apple, Disney, and IBM suspended their campaigns; the platform's ad revenue declined by more than half within a year (PLACEHOLDER: Insider Intelligence 2023 forecast; confirm exact citation). The episode was not isolated. The World Bank pulled advertising from X in 2024 after a paid placement appeared alongside content from a white supremacist organization (PLACEHOLDER: CBS/WSJ report citation). These incidents foreground a tension that confronts brands that advertise in social media: the same platforms that reach billions of consumers also expose branded content to environments brands cannot fully control. With global social media advertising spend exceeding $275 billion in 2025 (Statista 2025) and social channels attracting the largest share of incremental ad dollars worldwide (WARC 2025), the reputational and financial exposure and risk to brands might be substantial.

Marketing research examines this risk to brands through the lens of brand safety. Recent work in this tradition documents that the proximity of branded content to unsafe or negative material can contaminate brand perceptions, even when no explicit link between the brand and the adjacent content exists (Grewal, Vana, and Stephen 2025). The general recommendation, followed by both theory and practice, is that any association with this type of content should be avoided. Accordingly, brand safety strategies, for example detection algorithms or moderation, have become a major point of concern for both brands and platforms eager to avoid any unsafe content association. With regard to unsafe content, toxicity may represent a particularly potent source of such contamination, given its harmful and targeted nature. Yet, while brand safety concerns are real and demonstrably shape advertiser behavior, the actual reputational spillover may be weaker under naturalistic feed conditions than this experimental literature suggests. Prior designs typically isolate a single brand-content pairing, a setup that likely inflates contamination effects relative to the cluttered, high-velocity environment of an actual social media feed. Recent null findings are consistent with this concern (Bernritter et al. 2025).

Interestingly, a separate stream of research on platforms puts forward the idea that negative and toxic content can have positive implications (CITE). Negatively-valenced content, and toxic content specifically, captures disproportionate attention and generates higher engagement on digital platforms driven by algorithms and virality mechanisms (Beknazar-Yuzbashev et al. 2022;

ADD). The platform-level implication seems to be: toxic content retains and engages users. From this perspective, feeds with higher toxicity might expose co-located advertisements to larger and more engaged audiences.

These two streams generate an interesting context for the same managerial question: Should my brand withdraw from toxic discourse or not? Brand safety research advises withdrawal from toxic environments due to potential spillovers. Engagement research implies that those same environments could increase exposure, attention, and engagement. We argue that existing work cannot resolve this contradiction because each stream examines its mechanism in isolation. Many brand safety studies measure reputational or attitudinal outcomes without accounting for the attention environment that toxicity creates. Platform engagement studies measure platform-level metrics without tracking whether associative contamination offsets the attentional gains. Marketing research has yet to examine both within the same feed exposure, leaving brands and platforms without a basis for evaluating the net effect of feed-level and adjacent toxicity on advertising effectiveness.

We address this gap by integrating brand safety and platform engagement within a single empirical framework. Focusing on digital advertising on social media, where toxicity is widespread and brands exercise little control over adjacent content, we examine whether the relationship between feed-level toxicity and brand advertising effectiveness is more complex than either stream in isolation suggests. We manipulate the proportion of toxic content in simulated but more naturalistic social media feeds using a dose-response design (with stimulus sampling CITE) and measure both brand-directed (click-through rate, ad recall, brand attitude) and feed-level outcomes (dwell time, time spent on feed, user reactions).

Our central argument is that toxicity operates through two mechanisms. At moderate levels, it elevates feed-level engagement and thereby increases the probability that co-located ads receive attention, producing net-positive effects on engagement and attention. At high levels, the attentional benefit saturates and feed-level toxicity yields net-negative effects for the ad due to decreasing engagement with the feed (or a spillover at extreme levels?) Taken together, we expect the resulting relationship between feed-level toxicity and brand advertising effectiveness follows an inverted U. At the same time, direct reputational spillover from toxic content to the co-located brand is absent across most toxicity levels that are not extreme. Reputational spillover may be weaker under naturalistic feed conditions than prior experimental designs — which typically isolate a single brand-content pairing — would suggest, a possibility consistent with recent null findings (Bernritter et al. 2025). The engagement benefits documented in the opposing stream are real but bounded: they hold up to a threshold beyond which the feed environment turns hostile enough to suppress even the attentional gains.

Theory

Toxicity in Online Discourse

Toxicity describes intentionally inflammatory communication directed at a target to provoke unproductive conflict (Hanscom et al. 2024). It subsumes harassment, hate speech, cyberbullying, identity attacks, profanity, and threats, unified by a common function of inflicting harm on a target (Hanscom et al. 2024; Lees et al. 2022; Shahi and Majchrzak 2025). Crucially, toxicity is a property of communicative style rather than subject matter as it characterises how a message is deployed irrespective of the topic it addresses. Toxicity's purpose is to damage rather than to evaluate, criticize, or persuade per se. Constructive criticism, passionate disagreement, and even harsh evaluative language can be productive. Toxicity has no regard for disagreement or advancing an argument.

Because toxicity is defined by communicative function rather than topic, it follows that toxic discourse is pervasive across platforms, languages, and topics (Avalle et al. 2024; Shahi and Majchrzak 2025). It emerges wherever users interact, not only around highly controversial topics, such as politics, but also in entertainment, sports, and other consumer contexts (Salminen et al. 2021; Tufa, Markov, and Possen 2024). Platform architectures amplify this prevalence through recommendation algorithms, virality features, and network effects (Recuero 2024; Hanscom et al. 2024). Toxicity's consequences extend beyond the immediate online interaction. For instance, offline, inflammatory online content is linked to political unrest and hate crime (Müller and Schwarz 2021, 2023). For platforms, toxicity generates regulatory pressure: legislation in the EU and elsewhere mandates moderation of toxic content (Digital Services Act 2022; UK Online Safety Act 2023).

This definition also implies that toxicity is distinct from negativity as a theoretical construct.. Negativity is valence-defined: it encompasses any communication carrying negative sentiment, including product complaints, critical reviews, and expressions of dissatisfaction. These forms of negative communication can serve evaluative, corrective, or persuasive functions and are vital components of healthy discourse. Toxicity isolates the subset of negative communication that is specifically target-directed and harm-oriented. This distinction matters for marketing and advertising outcomes because the mechanisms differ. Negativity affects mood and may produce valence-based affective spillover (Bushman 2005; de Pelsmacker et al. 2002). Toxicity additionally activates threat detection and increases arousal that might heighten attention or even produce aversion at high levels (Öhman and Mineka 2001; Lenhart et al. 2016). An examination

grounded in negativity alone cannot account for these different mechanisms.

Table 1. The two-dimensional nature of toxic communication. Valence (positive/negative) and communicative style (civil/toxic) are orthogonal: toxic communication is defined by its harm-oriented function, not by its sentiment.

	Civil	Toxic
Positive	"Great game today, the team really pulled through in the second half."	"So proud of this squad. The fans rooting against us are pathetic and deserve to suffer."
Negative	"Disappointing result today. The defense looked completely disorganized and the midfield offered nothing."	"You're a disgrace to this club. Hope you break your leg next match, you useless idiot."

The same definition clarifies why toxicity differs from the content categories that govern programmatic brand safety practice and offers a complementary perspective for brand safety. Established industry standards or individual company solutions implement blocklists or exclusion checklists that classify content by topic: violence, adult material, hate speech, and similar categories (Amazon Ads, 2025; GARM, 2024[1]). Keyword blocklists implement these classifications by flagging topical terms regardless of communicative intent or intensity. These operationalizations are coarse and lead to type-1 and type-2 errors: On the one hand, keyword blocklists most likely fail to meet their goal whenever a discussion on a harmless topic turns toxic. A discussion on sports competitions, while generally considered a harmless topic, might turn toxic via its communicative style. On the other hand, these same blocklists might restrict advertising in harmless discussions for flagging keywords that are harmful only in specific contexts. These limitations are well documented. As an advertising executive noted: a "keyword blocklist approach frequently results in inventory that is contextually relevant, and safe, being blocked." For example, during the Euro 2024 final, the flagging of keywords, such as "shoot" or "attack" led to nearly half of a UK publisher's output being blocked from receiving advertising (The Media Leader 2024). In comparison, toxicity captures *how* users communicate, which is orthogonal to what these users communicate about. Relying on toxicity for brand safety might thus flag unsafe contexts with higher specificity (reducing over-blocking of safe inventory tied to type-1 errors) and precision (improving detection of toxic discussions on nominally harmless topics that topic-based approaches miss). At the same time, the captured dimension of communication style is what regularly produces attentional, attitudinal, and arousal effects relevant to advertising outcomes (see "Motivating Evidence" section).

[1]While GARM was discontinued in 2024, its underlying content categorization remains broadly in use (Meier, 2025).

Toxicity and Advertising

We examine two brand outcomes under feed-level toxicity. First, we evaluate ad effectiveness through (in)attention, which we approximate through unobtrusively measured behavioral measures (dwell time, clickthrough rate) and survey responses on memory (ad recall). Second, we investigate classic brand safety measures conceptualized as whether the advertising context harms the advertised brand. We measure these through survey responses on attitudinal indicators (brand attitude, brand trust). These outcome classes require separate theoretical treatment because they operate through different mechanisms.

Ad Effectiveness

Toxicity drives both attention and engagement. Negatively-valenced content in general captures disproportionate attention (Pratto and John 1991; Rozin and Royzman 2001). It receives higher engagement rates, greater sharing, and stronger recall than neutral and positive content (for an overview, see Robertson et al. 2023). Further, threat-relevant stimuli, such as toxic communication (e.g., personal insult), activate automatic detection processes that elevate attention (Öhman and Mineka 2001). Users who face the ad may process it shallowly or deeply depending on their arousal state at the moment of encounter. Excitation transfer theory holds that residual arousal from a preceding stimulus intensifies affective and cognitive responses to subsequent stimuli (Zillmann 1971). In a social media feed, toxic posts preceding the ad elevate the user's arousal through their provocative and targeted content. This residual arousal transfers to the ad. With toxicity, elevated arousal deepens processing of the ad.

Toxic content, specifically, has been demonstrated to generate larger, wider, and deeper reply threads than civil counterparts on Twitter (Saveski et al. 2021) and across other platforms (Avalle et al. 2024). The other way around, reducing toxicity on a platform via a browser extension showed a decrease in user attention and engagement across time spent, impressions, and clicks (Beknazar-Yuzbashev et al. 2025) in line with this reasoning. The provocative and targeted nature of toxic communication pulls in users who engage by replying, retaliating, defending, or spectating. Overall, toxicity likely increases overall attention on feed elements, including the ad.

These benefits of toxicity likely come to a halt at high toxicity. Sustained exposure to severely toxic content like

hate speech, harassment, dehumanizing language, produces aversion, lowered engagement, and withdrawal. Users report that they avoid participation for fear of harassment or shut down their accounts entirely; self-censorship and silencing are among the observed reactions (Lenhart et al. 2016). For instance, misogynistic hate speech reduces women's participation in political discourse (Carlson 2018). In gaming, toxicity decreases player engagement (Liu et al. 2025). Put differently, studies on content moderation find that reducing toxicity can increase participation (Matias 2019; Wise et al. 2006). Furthermore, the operationalization of toxicity that became the industry standard (Aridor et al. 2024) directly includes a heightened propensity to leave a discussion as a signal for its presence (Lees et al. 2022).

The traditional cue-utilization hypothesis is informative as well as it predicts a reversal in processing depth at high levels of arousal. Elevated arousal from negative affect narrows the range of environmental cues that receive processing resources (Easterbrook 1959). Under moderate toxicity, the attentional window still encompasses both toxic content and adjacent stimuli. Under high toxicity and thus arousal, attention concentrates on threat-relevant cues at the expense of peripheral stimuli (Martha and Sutherland, 2014; van Steenbergen et al. 2011). Users in a state of heightened threat vigilance prioritize monitoring toxic content and disengage from elements perceived as irrelevant to the arousal source. The ad, a commercially oriented, non-threatening element, thus loses attentional resources.

Ad effectiveness therefore depends on the overall levels of toxicity in the feed. At low-to-moderate doses, arousal likely exceeds aversion. At high doses, the relationship reverses. We therefore expect an inverted U: engagement increases with toxicity dose up to a threshold and declines beyond it.

> **H1:** Feed-level toxicity and ad engagement (ad dwell time, ad clickthrough rate, ad recall) follow an inverted-U relationship: ad engagement increases with the proportion of toxic content up to a moderate threshold and decreases at higher proportions.

Brand Safety

Spillover contamination from adjacent content is a central mechanism in brand safety research: consumers transfer negative associations from proximate content to a co-located brand, even absent any causal or semantic link between them (Grewal, Vana, and Stephen 2025). That is, when a brand is presented alongside unsafe content online, there is a spillover of the adjacent negativity that harms said brand. This spillover of contextual factors on co-located brands presented in the same context has been demonstrated in other settings, such as print and television (Bushman 2005; de Pelsmacker et al. 2002; Nian et al. 2021). Because toxicity is defined by its harm-oriented function, it constitutes unsafe content by definition, and spillover theory therefore predicts that toxicity in the content surrounding an ad harms brand attitudes.

The attention mechanism underlying H1, however, introduces level-dependence that prior spillover theory does not anticipate. Contamination requires encoding of the brand-context association, and encoding depth is a function of attentional resources allocated to the ad. At moderate toxicity (where H1 predicts the highest attention allocation) deeper processing of the ad occurs within a hostile environment, which should strengthen the brand-context association and amplify contamination. At high levels toxicity (where H1 predicts the withdrawal attentional resources toward threat-relevant cues) the ad receives fewer processing resources. This would weaken the association through which contamination operates even as the environment becomes more hostile. The net effect on brand attitudes is thus expected to resemble the attentional pattern in H1: contamination should be strongest at moderate toxicity and attenuate at high toxicity.

> **H2:** Toxicity in the content surrounding an ad harms brand attitudes, with the effect being strongest at moderate toxicity levels and attenuating at high toxicity.

Whether H2 holds in a naturalistic feed context is an open empirical question. Prior evidence for spillover contamination predominantly stems from studies that isolate brand-content pairings. For example, Lee et al. (2021, e.g., Study 1) presented participants with a single offensive social media post on animal abuse with a pet shop brand ad presented juxtaposed to the post. While such designs are certainly powerful to manipulate and capture psychological processes of brand safety concerns, they tend to neglect the feed structure, that is, the actual browsing in which consumers encounter brand safety incidents. This may affect the predicted effects in two ways. First, these studies concentrate processing resources on the brand-content juxtaposition and are likely to induce an overly analytical mindset [Pham (2013); Sawyer 1975]. This suggests that prior estimates may overstate what occurs under feed conditions. Second, prior designs operationalize local adjacency rather than feed-level toxicity. What typically varies is the content directly neighboring the ad, not the toxicity of the entire feed. In contrast, the present study varies feed-level toxicity independently of adjacent-post toxicity. Doing so, we intend to estimate a quantity that prior research has not directly addressed. How strongly spillover contamination operates through feed-level toxicity, therefore, remains an open question.

Motivating Evidence

Research on social media platforms shows that toxic content drives *engagement*: it attracts more reactions, shares, and time spent than civil alternatives (Beknazar-Yuzbashev et al. 2022; Aridor et al. 2024; Beknazar-Yuzbashev, Jiménez-Durán, and Stalinski 2024). We provide evidence that toxicity also drives *attention* specifically, using 12,473 A/B tests from the *Upworthy Research Archive* (Matias et al. 2021) that cover 53,755 news headlines that were displayed 205.8M times.

In this setting, each A/B test randomly assigned website visitors to one of several headlines referring to the same article and recorded clicks on each headline. Unlike likes or comments, these 2.8M clicks were not visible to others. The clicks in this dataset therefore reflect attentional responses rather than social behavior, which is precisely the channel our theoretical argument concerns.

The recorded click-through rates (CTRs) range from 0.000% to 14.891%. The average CTR across all headlines is 1.394% and the median click rate is 1.067%. As the distribution among CTRs is right-skewed, only a small proportion of news stories are associated with a high CTR.

Toxic news headlines, defined as those exceeding a Perspective API score of 0.6, achieve higher click-through rates than civil alternatives ($Median_{CTR}^{Civil} = 1.061\%$; $Median_{CTR}^{Toxic} = 1.401\%$). Figure 1 provides model-free support for this claim: the distribution of click-through rates is shifted rightward for toxic headlines across the full range of observed values. A fixed-effects regression that compares headlines which refer to the same article quantifies the difference: the estimated uplift corresponds to 6.408% of the median click-through rate, and it holds after controlling for headline length, readability, and negative emotional valence.

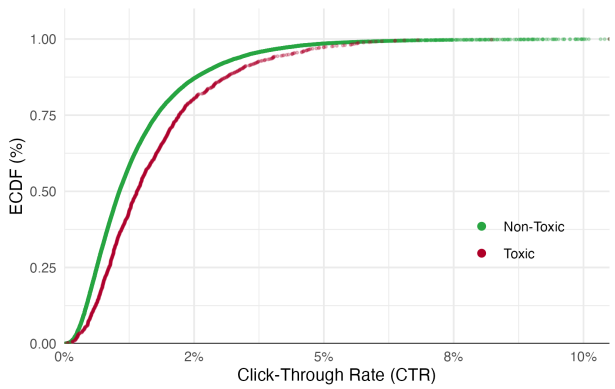

Figure 1. Empirical cumulative distribution functions of toxic vs. non-toxic headlines

These results suggest that toxicity drives attention in a setting where article content is held constant and social motives are absent. Importantly, the Upworthy data represent an unusually clean source of field evidence for this purpose as the A/B tests were designed by the platform itself: visitors were randomly assigned to headline variants, one experiment ran at a time, and commercial incentives ensured careful implementation. Such data are rare. Even so, three limitations prevail. First, the data contain no information about the feed structure in which each headline was embedded, which is why we cannot derive a dose-response relationship between toxicity and attention. Furthermore, this study operationalizes clicks on the stimulus itself. It does therefore not offer evidence of excitation transfer from toxic content to adjacent, non-toxic content. Finally, it does not provide evidence on downstream consequences such as brand attitudes.

Addressing these limitations with additional field data that is typically available to marketing scholars is difficult, if not impossible. This is because online platform studies and observational studies both suffer from algorithmic interference: platform algorithms determine which users see which content in ways that are unobservable to researchers, making it difficult to isolate the effect of feed composition on any outcome (Xu, Zhang, and Zhou 2020; Johnson 2023; Braun and Schwartz 2025). Furthermore, neither paradigm allows researchers to manipulate or even observe the context surrounding a focal ad, which is required to study brand safety (Roggenkamp, Boegershausen, and Hildebrand 2025). In what follows, we therefore turn our attention to an experimental design.

Experimental Evidence

To address the three limitations of the news headline study, we designed a Digital In-Context Experiment to embed a sponsored post within a feed whose toxicity dose is experimentally varied. This enables a clean operationalization as well as more ecologically valid stimuli and measures compared to typical vignette experiments (Roggenkamp, Boegershausen, and Hildebrand 2025).

[Give a brief overview here]

Scenario (WIP, Claude's Version)

Participants are exposed to a scenario around the movie *Avatar: Fire and Ash* (2025). Topic selection for a feed toxicity study requires a balance between ecological validity and experimental control. The topic must generate genuine toxic discourse in sufficient volume to construct a large stimulus corpus, engage participants authentically, and avoid partisan or identity-based reactions strong enough to saturate outcome measures independently of the toxicity manipulation. Entertainment content about a major release satisfies these criteria. Film discourse is emotionally engaging — participants form genuine opinions and reactions — yet is not structured by political or identity cleavages that would produce floor or ceiling effects on engagement or attitudinal outcomes. The topic also constitutes a realistic context for consumer advertis-

ing, which supports the ecological validity of the brand safety design.

Design

The study employs a dose-response design, implemented using the Digital In-Context Experiment (DICE) paradigm (Roggenkamp, Boegershausen, and Hildebrand 2025; Chen, Schonger, and Wickens 2016). We expose participants to a social media feed in which we manipulate the proportion of toxic content exogenously and continuously between participants. DICE renders the feed in a responsive browser interface that replicates the visual structure of a microblogging platform and records granular behavioral data unobtrusively and without requiring any explicit response from participants. Each participant views a single feed consisting of 25 organic posts and one sponsored post. The toxicity dose d, defined as the proportion of toxic organic posts, ranges from 0 (a fully civil feed) to 1 (a fully toxic feed) in increments of $\frac{1}{25}$. This continuous manipulation allows estimation of the full dose-response function, which is a prerequisite for testing the inverted-U predictions in H1 and H2.

Constructing valid stimuli for a toxicity manipulation (e.g., via web scraping) is complicated by the fact that toxicity often correlates with several other content attributes. Toxic posts tend to be more profane, more negatively valenced, more emotionally intense, and more controversial than civil posts. Manually retrieving and constructing matched pairs that differ only in toxicity is therefore difficult (if not infeasible), because any revision that removes toxic elements changes these correlated dimensions as well. We address this issue by combining two methodological approaches: stimulus sampling (Simonsohn, Montealegre, and Evangelidis 2025) and scope testing with AI-generated stimuli (Tomaino et al. 2025). Accordingly, we treat individual posts as a random factor and draw each participant's feed independently from a large, pre-constructed stimulus corpus. Because idiosyncratic features of specific posts vary randomly across participants assigned the same dose, potential confounds are uncorrelated with the dose manipulation in expectation. In addition, we use generative AI to construct the corpus itself. This makes it tractable to generate posts across the full toxicity range while systematically varying orthogonal dimensions such as length, sentiment polarity, writing style, and thematic focus.

The stimulus corpus was constructed in three stages. In the first stage, we prompted the Grok-3 API (xAI) to generate posts across four broad toxicity levels (i.e., civil, mildly toxic, toxic, and highly toxic) while independently varying length, writing style, sentiment, and thematic focus. Each post was scored using the Google Perspective API, which provides a continuous toxicity score on the $[0, 1]$ scale and constitutes the industry standard for computational toxicity detection in research and practice (Aridor et al. 2024). The initial stimuli generation produced adequate coverage at the extremes of the distribution but left a gap in the moderate range (scores between approximately 0.15 and 0.65). A second stage targeted this gap. Specifically, we applied a generate-and-filter procedure: posts were generated to three intermediate toxicity targets and accepted only if their toxicity score (predicted with the Perspective API) fell within the target band, otherwise generation was repeated. In a third stage, sampled posts were rephrased using another Grok-3 API call to eliminate repetitive phrasing patterns. The resulting corpus (consisting of 461 organic posts) almost spans the full toxicity range [0.019, 0.945] with near-uniform coverage.

Posts in the corpus are classified as toxic (Perspective score ≥ 0.6) or non-toxic (i.e., civil; score < 0.6), which yields two sampling pools. For a participant assigned dose d, the feed is assembled by sampling $d \times 25$ posts without replacement from the toxic pool and $(1 - d) \times 25$ posts from the civil pool, then randomly shuffling the resulting 25 organic posts. Each participant draws an independent sample, so participants assigned the same dose see different posts in a different order.

The sponsored post advertises M&M's and invites participants to enter a prize lottery by clicking on the post. Participants who click are entered into a real lottery. The brand and product category were chosen to match the scenario framing: a snack brand is a contextually plausible advertiser in a feed about a movie. The sponsored post is inserted at a fixed position (14th scroll position) and is identical across all participants and dose conditions.

Procedure

After providing informed consent, participants read a brief scenario framing the task: they had recently come across *Avatar: Fire and Ash* and were opening a microblogging feed to see what others were saying about the film. They are instructed to browse at whatever pace and depth feels natural and to click "Done browsing" when they feel they have seen enough. DICE records behavioral data unobtrusively throughout: the viewport tracker logs dwell time continuously for each post as it enters and exits the visible area, and participants may like posts or compose reply text within the feed.

After clicking "Done browsing," participants are redirected automatically to a Qualtrics survey. To preserve the integrity of the unaided recall measure, the survey opens with the recall questions before participants are asked about the brand directly. Unaided recall is elicited first, followed by aided recall, and the brand attitude battery closes the survey along with a prior brand knowledge item.

Outcomes

Participant behavior during feed exposure is recorded unobtrusively. A viewport tracker logs the duration each

post is visible in the browser window (dwell time) to provide a proxy for in(attention). Participants could also click on the ad (to sign up for an incentive compatible lottery), like posts and compose reply text, which is recorded as indicators of active engagement. Total time from loading the feed in the browser to proceeding to the next page was recorded as overall time on feed. We also record paradata, such as the device type and screen resolution to allow for covariate adjustment for differences in feed rendering across devices.

Outcomes are ordered by their theoretical proximity to the attention mechanism and their statistical sensitivity, establishing a pre-registered primary/secondary hierarchy to control the family-wise Type I error rate.

Ad dwell time

This (log-transformed) variable is the primary outcome. It is the most direct unobtrusive measure of attention allocated to the sponsored post, recorded continuously by the viewport tracker without requiring any active response from participants. Viewport dwell time shows high concurrent validity with mobile eye-tracking measures ($r = .89$ with gaze duration) and accounts for 76% of the variance in brand recall in validation data (Bruns et al. 2025). The sample size is powered for this outcome. H1 is supported or rejected on the basis of this outcome alone. Something we should check in our pilot data, too.

Recall

Unaided recall is a secondary outcome. Participants are asked to list any brands from which they recall having seen an ad in the feed. This captures whether exposure was sufficient to encode the brand into memory without a cue, and is stricter and less susceptible to guessing than aided recall. Aided recall, in which participants select from a list of eight brands the one they recall seeing a sponsored post for, is reported as a robustness check when unaided recall is sparse.

Lottery signup

This is binary variable. which we also consider as is a secondary outcome. Although behavioral and also a good approximation of financial brand safety (Bernritter et al. 2025), it is relatively sparse, which limits power. It is therefore reported for completeness but not treated as confirmatory.

Time on feed

Similar to the dwell time measure, this log-transformed outcome is reported as a feed-level variable. It addresses whether feed-level toxicity affects overall browsing duration. As such, the variable may predict (in)attention, but is not an ad effectiveness measure.

Brand attitude

is the primary outcome for H2. Participants rate agreement with six statements about M&M's on a 7-point scale (1 = *strongly disagree*, 7 = *strongly agree*); items are drawn from Grewal, Vana, and Stephen (2025) and cover purchase consideration, quality perception, recommendation likelihood, social identification, general liking, and trustworthiness. The composite score is the mean across all six items. We apply the same two-lines procedure and decision rules as for H1 (see Table 2). If neither the two-lines test nor the quadratic term is significant, we conclude that feed-level toxicity does not affect brand attitude detectably.

Pilot Data

We recruited 201 participants for a pilot study. We operationalized feed toxicity differently for the pilot, compared to how we will operationalize it in the future. Specifically, we used a Gaussian kernel centered on target mean scores drawn from a uniform grid. Realized feed toxicity ranged from 0.105 to 0.842. The main study replaces this operationalization with the binary proportion method described in the Design section, which assigns toxicity dose as a fixed proportion of toxic posts and ensures uniform coverage across the full range.

The sponsored post entered the viewport of 136 participants (about two thirds of all participants). The remaining 65 did not scroll far enough for the ad to become visible. Among the 136 exposed participants, median ad dwell time was 1.936 seconds (mean 3.348s, SD 4.791s). On the log(1+s) scale used in the primary analysis, the pilot mean is 1.164 and SD is 0.716. These estimates ground the power analysis. Four participants spent fewer than 20 seconds browsing the feed.

Recall rates indicate adequate detection sensitivity. Seventy-two of 201 participants (35.8%) recalled M&M's on the unaided measure; 94 (46.8%) identified it on the aided measure. The gap between unaided and aided recall supports that both measures are informative and justifies treating aided recall as a robustness check. Lottery signup was sparse (6/201, 2.985%), consistent with the expectation in the Outcomes section. Brand attitude showed a pilot mean of 5.235 (SD 1.376) on the 1-7 scale, with no evidence of a ceiling effect.

Analysis Plan

H1: Feed toxicity and ad engagement

H1 predicts that feed-level toxicity and ad engagement follow an inverted-U relationship.

All secondary outcomes are tested at $\alpha = .05$ using the same two-lines procedure as the primary outcome but are not the basis for supporting or rejecting H1. [We should rewrite this. This is intended to be a disclaimer that we will not correct for multiple hypothesis testing because the outcomes are secondary and because we only have on primary outcome per hypothesis (which are more or less unrelated).]

We follow the two-lines procedure proposed by Simonsohn (2018). For each outcome, we 1. fit a cubic spline to estimate the location of the peak, 2. apply the Robin Hood algorithm to determine the break point that maximizes

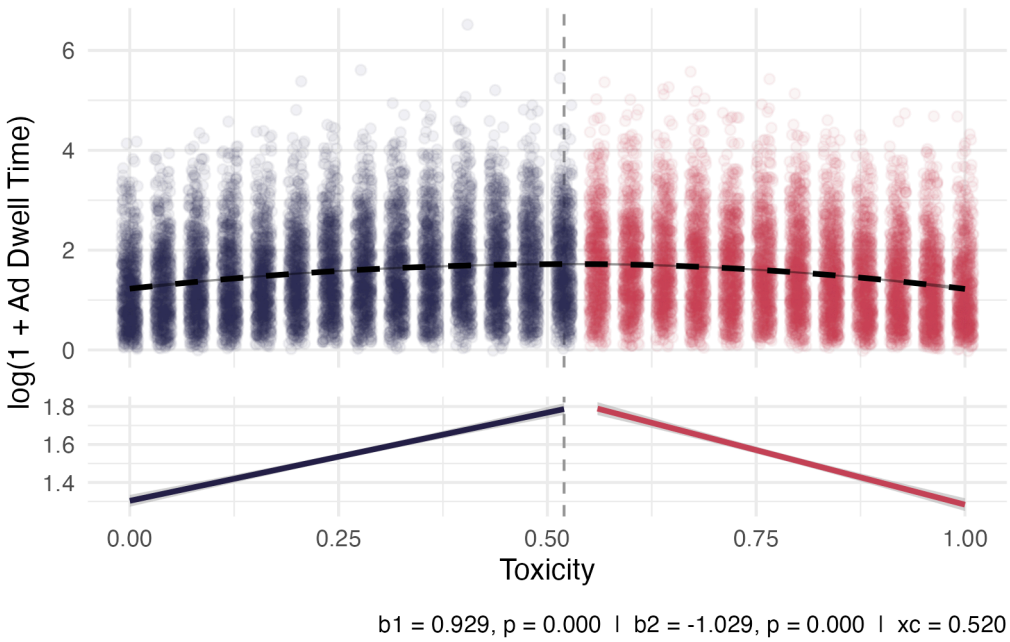

Figure 2. Two-lines test: feed toxicity and ad dwell time. Upper panel shows raw data with LOESS smoother; lower panel shows the two fitted regression lines. Break point determined by the Robin Hood procedure (Simonsohn 2018).

power by allocating observations toward the weaker of the two slopes, and 3. estimate an interrupted regression with one slope below and one slope above the break point. H1 is fully supported if both slopes (denoted b_1 and b_2 in Table 2) pass one-sided tests ($\alpha = .05$) in the predicted directions: $b_1 > 0$ below the break point and $b_2 < 0$ above it. Because the two-lines test requires a true sign reversal, a positive-then-flat trajectory — in which engagement increases and then saturates without declining — would not satisfy the strict criterion but is still broadly consistent with the arousal mechanism and is treated as partial support.

To capture this possibility, we additionally estimate a quadratic regression of the outcome on `feed_toxicity` and `feed_toxicity`^{2}. A significant negative quadratic coefficient alongside the plateau pattern in the two-lines test is interpreted as partial support for H1: the data show the predicted initial increase and a curvilinear compression, even if the second slope does not turn significantly negative. This secondary test is reported for completeness but, following Simonsohn (2018), is not treated as primary evidence given the known tendency of quadratic regression to produce false positives. The full decision rules are in Table 2.

Table 2. Pre-registered decision rules for H1. Quadratic term column is only evaluated when the two-lines test does not yield opposite significant slopes.

Two-lines result	Quadratic term	Interpretation
$b_1 > 0^*, b_2 < 0^*$	—	**Support H1** (inverted-U supported)
$b_1 > 0^*, b_2$ n.s.	$< 0^*$	**Partial support H1** (plateau: engagement increases and saturates but does not invert)
$b_1 > 0^*, b_2$ n.s.	n.s.	Reject H1 (monotonic increase)
$b_1 > 0^*, b_2 > 0^*$	—	Reject H1 (monotonic increase)
b_1 n.s. or b_2 n.s.	—	Reject H1 (no detectable relationship)

H2: Feed toxicity and brand attitude

H2 predicts an inverted-U relationship between feed-level toxicity and brand attitude, with contamination

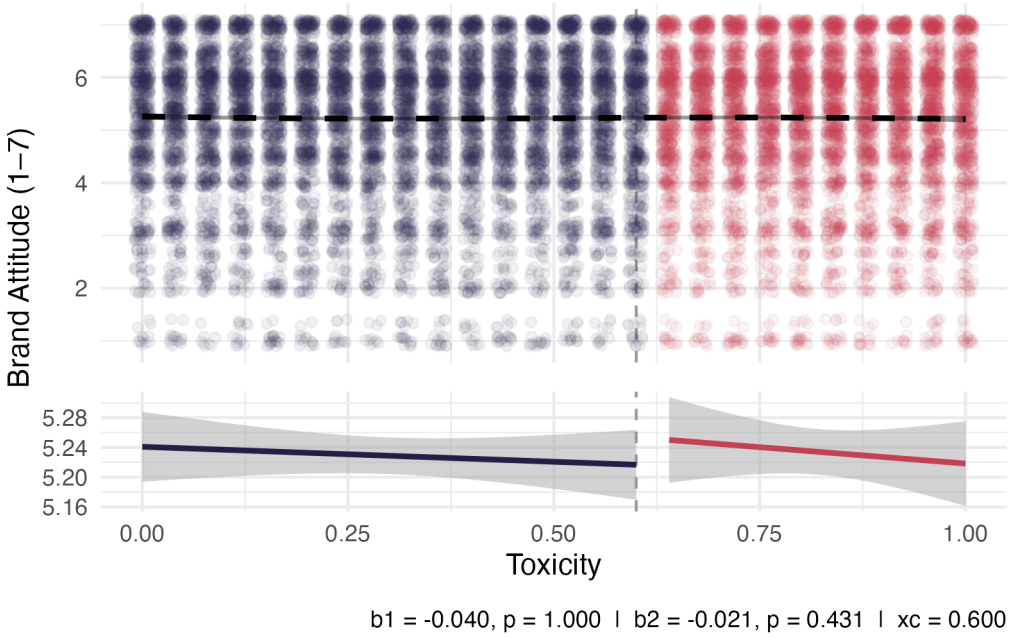

Figure 3. Two-lines test: feed toxicity and brand attitude.

strongest at moderate toxicity and attenuating at high toxicity. We apply the same two-lines procedure and quadratic fallback as for H1 (see Table 2). The primary test is whether both slopes are significant with opposite signs; partial support is declared if the first slope is positive and significant and the quadratic term is significantly negative.

We compute a composite brand attitude score as the mean of the six focal brand items (`gfb_1` to `gfb_6`). If neither the two-lines test nor the quadratic term is significant, we conclude that feed-level toxicity does not affect brand attitude in a detectable way, consistent with the attenuation argument in the brand safety section above.

Robustness and Potential Concerns

Because organic post order is randomized within each feed, the posts appearing immediately above and below the sponsored post vary across participants. The local toxicity of the ad's immediate context therefore co-varies with the feed-level dose in expectation but is not independently manipulated. To assess whether the estimated dose-response relationship is driven by local adjacency rather than overall feed composition, we re-estimate the primary analysis controlling for the mean Perspective API score of the two organic posts adjacent to the sponsored post (one above, one below). If the dose effect is attenuated once local adjacency is held constant, this would suggest that proximity to toxic content — rather than ambient feed-level toxicity — is the operative mechanism. Stability of the estimate would support the feed-level interpretation.

Sample Size Justification

The two-lines test has no closed-form power formula, so we estimate power by simulation following Lakens (2022). For each sample size $N \in [200, 400, \ldots, 4800, 5000]$, we

1. generate a dataset with two assumed inverted-U data-generating processes representing for small and medium effect sizes,
2. apply the two-lines test,
3. flag a success when both slopes pass a one-sided test ($\alpha = .05$; $b_1 > 0$ and $b_2 < 0$), and
4. repeat the procedure 1,000 times.

Power is defined as the probability that the test correctly rejects the null hypothesis when there actually is a genuine effect. In our simulation, power is therefore represented as the proportion of successful runs.

A key concept throughout is the inverted-U peak: under the assumed model $\mathbb{E}[Y \mid x] = \beta_0 + \beta_1 \cdot x(1-x)$, the function $x(1-x)$ reaches its maximum of 0.25 at $x = 0.5$ (a feed that is exactly half toxic). The peak outcome is therefore $\beta_0 + 0.25\,\beta_1$, and $0.25\,\beta_1$ is the peak uplift — how much higher the expected outcome is at 50% toxicity compared to a fully clean or fully toxic feed. Effect sizes are expressed as Cohen's d at this peak: $d = 0.25\,\beta_1/\sigma$, or equivalently $\beta_1 = 4\,d\,\sigma$.

We report power separately for two outcomes (ad dwell time and brand attitude) using this common d scale. The power curves (Figure 4) indicate that detecting a small effect ($d = 0.2$) with 80% power requiresapproximately $N = 4,000$ participants; we adopt this as our recruitment target. The two-lines test is more conservative than a

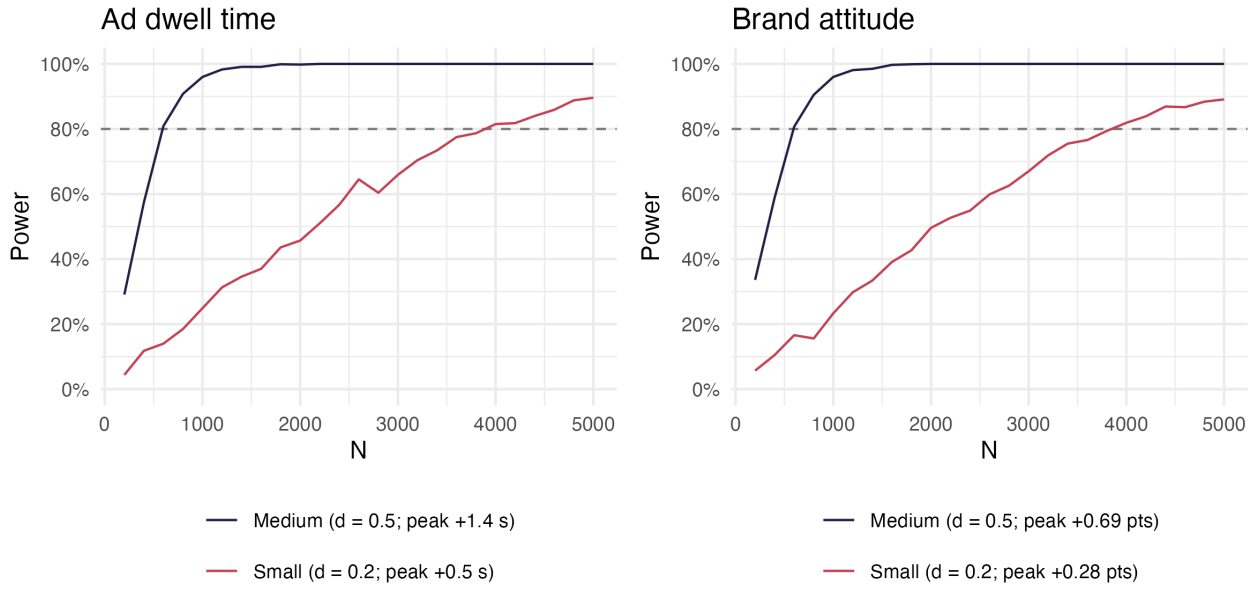

Figure 4. Power curves for the two-lines test on ad dwell time (left) and brand attitude (right). Effect sizes are expressed as Cohen's d at the inverted-U peak (d = peak uplift / pilot SD). Dashed line marks 80% power.

Recruitment

Exclusion Criteria

simple quadratic regression, so this sample size provides at minimum adequate power for the pre-registered primary test.

Ad dwell time. Because ad dwell time is heavily right-skewed, we analyze it on the $\log(1 + \text{seconds})$ scale, with σ estimated from the pilot. The pilot baseline is 2.2 s (median dwell time at the tails, i.e., in a fully clean or fully toxic feed). The peak — a 50% toxic feed — raises median dwell time to 3.6 s under a medium effect ($d = 0.5$), an absolute increase of 1.4 s. Under a small effect ($d = 0.2$), the peak median is 2.7 s.

Brand attitude. We apply the same d = 0.5 / d = 0.2 convention. The pilot mean is 5.24 on a 1–7 scale (SD: 1.38). A medium effect raises peak-condition attitude by 0.69 points; a small effect by 0.28 points.

References

Aridor, Guy, Rafael Jiménez-Durán, Ro'ee Levy, and Lena Song. 2024. "The Economics of Social Media." *Journal of Economic Literature* 62 (4): 1422–74. https://doi.org/10.1257/jel.20241743.

Barberá, Pablo, Amber E. Boydstun, Suzanna Linn, Ryan McMahon, and Jonathan Nagler. 2020. "Automated Text Classification of News Articles: A Practical Guide." *Political Analysis* 29 (1): 19–42. https://doi.org/10.1017/pan.2020.8.

Beknazar-Yuzbashev, George, Rafael Jiménez Durán, Jesse McCrosky, and Mateusz Stalinski. 2022. "Toxic Content and User Engagement on Social Media: Evidence from a Field Experiment." Working Paper, November. https://doi.org/10.2139/ssrn.4307346.

Beknazar-Yuzbashev, George, Rafael Jiménez-Durán, and Mateusz Stalinski. 2024. "A Model of Harmful yet Engaging Content on Social Media." *AEA Papers and Proceedings* 114 (May): 678–83. https://doi.org/10.1257/pandp.20241004.

Bernritter, Stefan F., Ilias Danatzis, Jana Möller-Herm, and Francesca Sotgiu. 2025. "Navigating Toxic Playgrounds: Managing Reputational and Financial Brand Safety in Multiplayer Video Games." *International Journal of Research in Marketing* 42 (4, Part A): 936–56. https://doi.org/10.1016/j.ijresmar.2025.01.002.

Braun, Michael, and Eric M. Schwartz. 2025. "Where a/b Testing Goes Wrong: How Divergent Delivery Affects What Online Experiments Cannot (and Can) Tell You about How Customers Respond to Advertising." *Journal of Marketing* 89 (2): 71–95. https://doi.org/10.1177/00222429241275886.

Bruns, Daniel, Julian Felix Kopka, Lennart Borgmann, Steffen Prior, and Tobias Langner. 2025. "Measuring Gaining and Holding Attention to Social Media Ads with Viewport Logging: A Validation Study Using Mobile Eye-Tracking." *Journal of Advertising* 54 (5): 655–72. https://doi.org/10.1080/00913367.2025.2524186.

Chen, Daniel L., Martin Schonger, and Chris Wickens. 2016. "oTree—an Open-Source Platform for Laboratory, Online, and Field Experiments." *Journal of Behavioral and Experimental Finance* 9: 88–97. https://doi.org/10.1016/j.jbef.2015.12.001.

Grewal, Lauren S., Prasad Vana, and Andrew T. Stephen. 2025. "Brands in Unsafe Places: Effects of Brand Safety Incidents on Brand Outcomes." *Journal of Marketing Research* 62 (6): 981–1002. https://doi.org/10.1177/00222437251349522.

Johnson, Garrett A. 2023. "Inferno: A Guide to Field Experiments in Online Display Advertising." Journal Article. *Journal of Economics & Management Strategy* 32 (3): 469–90. https://doi.org/10.1111/jems.12513.

Lakens, Daniël. 2022. "Improving Your Statistical Inferences." Zenodo. https://doi.org/10.5281/zenodo.6409077.

Matias, J. Nathan, Kevin Munger, Marianne Aubin Le Quere, and Charles Ebersole. 2021. "The Upworthy Research Archive, a Time Series of 32,487 Experiments in u.s. Media." *Scientific Data* 8: 195. https://doi.org/10.1038/s41597-021-00934-7.

Pham, Michel Tuan. 2013. "The Seven Sins of Consumer Psychology." *Journal of Consumer Psychology* 23 (4): 411–23. https://doi.org/10.1016/j.jcps.2013.07.004.

Roggenkamp, Hauke, Johannes Boegershausen, and Christian Hildebrand. 2025. "DICE: Advancing Social Media Research Through Digital in-Context Experiments." *Journal of Marketing*. https://doi.org/10.1177/00222429251371702.

Simonsohn, Uri. 2018. "Two Lines: A Valid Alternative to the Invalid Testing of u-Shaped Relationships with Quadratic Regressions." *Advances in Methods and Practices in Psychological Science* 1 (4): 538–55. https://doi.org/10.1177/2515245918805755.

Simonsohn, Uri, Andres Montealegre, and Ioannis Evangelidis. 2025. "Stimulus Sampling Reimagined: Designing Experiments with Mix-and-Match, Analyzing Results with Stimulus Plots." *Journal of Personality and Social Psychology* 129 (1): 71–90. https://doi.org/10.1037/pspa0000449.

Tomaino, Geoff, Asaf Mazar, Ziv Carmon, and Klaus Wertenbroch. 2025. "A Simple Method for Improving Generalizability in Behavioral Science: Scope Testing with AI-Generated Stimuli (STAGS)." *Consumer Psychology Review* 8 (1): 87–97. https://doi.org/https://doi.org/10.1002/arcp.1101.

Xu, Heng, Nan Zhang, and Le Zhou. 2020. "Validity Concerns in Research Using Organic Data." *Journal of Management* 46 (7): 1257–74. https://doi.org/10.1177/0149206319862027.

Appendix

Stimulus Corpus Construction

Rationale for AI-Generated Stimuli

Constructing a stimulus corpus that spans the full toxicity range while holding other content attributes constant is not feasible with scraped content. Toxic posts in the wild tend to be shorter, more negatively valenced, more emotionally intense, and more syntactically informal than civil posts. Manually matching or adjusting scraped posts to isolate toxicity from these correlated attributes is both labor-intensive and practically limited: any manual revision that removes toxic language also changes tone, length, and affect. We therefore generate the corpus synthetically using a large language model (see also Tomaino et al. 2025), which allows systematic variation of toxicity level as an independent dimension while orthogonally varying post length, writing style, thematic focus, and format. Something we should test.

Concerns about ecological validity are addressed through two design choices. First, we calibrate all generated posts against the Google Perspective API (Aridor et al. 2024), the industry-standard classifier used in platforms' own moderation systems, which provides an objective anchor independent of the generation process. Second, we rephrase all posts through a separate humanization step that eliminates repetitive LLM phrasing patterns and introduces realistic orthographic variation (e.g., missing apostrophes, dropped word-endings) calibrated to the level present in organic social media text. The humanization prompt is anchored by a set of real posts scraped from the social media platform X about *Avatar: Fire and Ash*, which serve as few-shot register examples. These design choices are intended to ensure that the posts read as authentic organic discourse rather than machine-generated text, while preserving the controlled toxicity distribution that natural data cannot provide.

Generation Pipeline

The corpus was constructed in four stages: (1) initial generation across the full toxicity spectrum, (2) gap-filling via a generate-and-filter loop, (3) humanization and toxicity rescoring, and (4) AI-based content annotation.

Stage 1: Initial generation. We prompted the Grok-3 API (xAI) to generate posts about *Avatar: Fire and Ash* (2025) across four broad toxicity levels: civil, mildly toxic, toxic, and highly toxic. Each generation call independently sampled from a set of orthogonal variation dimensions – post length (very short to long), writing style (casual, formal, emphatic, abbreviated, deadpan, meme-register), thematic focus (specific characters, box office performance, visual effects, personal reactions, comparisons to earlier entries in the franchise), and post format (statement, rhetorical question, reaction, listicle, fan theory). A factual context block injected into the system message grounded all posts in verified information about the film (cast, characters, box office gross, critical reception, Cameron's stated plans for sequels). A discourse-hook dimension sampled from a list of recurring fan debate topics (e.g., the underuse of the villain Varang, the divisive portrayal of Spider, Disney's reported plans to reduce future budgets) to ensure posts reflected the actual topical structure of audience discourse rather than generic movie reactions.

Stage 2: Gap-filling. Initial generation produced adequate coverage at the low and high extremes but left sparse coverage in the moderate range (Perspective scores approximately 0.15-0.65) and in the extreme upper tail (scores above 0.80). Two supplementary scripts addressed these gaps using a generate-and-filter loop: posts were generated with target-band guidance and accepted only when their Perspective API score fell within the specified interval; otherwise generation was retried up to a fixed maximum. The moderate-range script targeted three sub-bands (0.15-0.35, 0.35-0.50, 0.50-0.65). The extreme-range script crossed two upper bands (0.65-0.80, 0.80-1.00) with five content moderation categories (harassment, hate speech, threatening language, sexually explicit content, and content promoting illicit activity, following the OpenAI moderation taxonomy), ensuring variety within the toxic pool beyond simple profanity. A fourth script targeted the low end of the civil range (0.00-0.15), covering fan enthusiasm, craft appreciation, and emotional responses that are absent from generation pipelines optimized for toxicity variation. All generation scripts used the same factual context block, discourse hooks, and variation dimensions as Stage 1.

Stage 3: Humanization and rescoring. All generated posts were pooled and sampled to approximate a uniform distribution across ten equal-width toxicity bins (target: 50 posts per bin; 500 posts total). Because humanization systematically shifts Perspective scores downward (rephrased posts were generally less aggressive in surface form than the originals) the three highest bins were oversampled before humanization (by factors of 2, 3, and 5 respectively) and re-capped at the target after rescoring. Sampled posts were sent in batches to Grok-3 with instructions to rephrase each post so it reads as a different real person's contribution, while preserving the toxicity intensity indicated by its Perspective score. The rephrasing prompt included a structural avoidance list (overused openings and trigrams detected by n-gram frequency analysis across the sampled pool), a set of nine register-calibrated few-shot examples spanning the full toxicity range, and instructions varying post length, movie reference style, and thematic angle. All posts were then rescored with the Perspective API, and the final corpus was assembled by re-capping at 50 posts per bin on the rescored distribution. Typos were injected post-hoc.

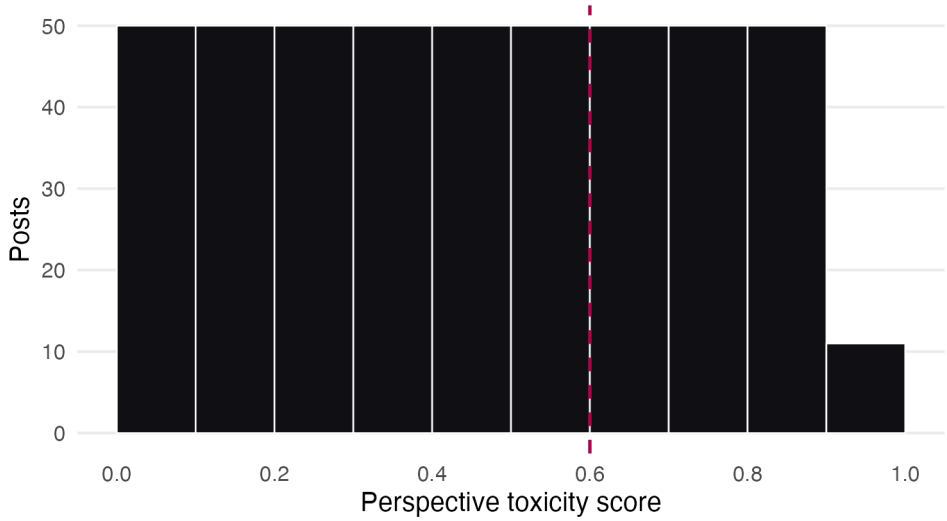

Figure 5. Distribution of Perspective API toxicity scores in the stimulus corpus. Dashed vertical line marks the conventional toxic/civil threshold (0.6). Bars correspond to the ten sampling bins used to ensure coverage across the full toxicity range.

Stage 4: Content annotation. Each post in the final corpus was annotated using two complementary AI-based scoring systems. First, the OpenAI `omni-moderation-latest` model provided category-level probability scores across eleven content dimensions (e.g., harassment, hate speech, sexual content, violence), characterizing the content type of each post independently of its Perspective toxicity score. Second, GPT-4o-mini (version `gpt-4o-mini-2024-07-18`) provided five ratings per post via structured JSON output: a continuous toxicity score (0–1) using a definition that does not reference Perspective API:

> *"language that a reasonable person would find attacking, devaluing, or harmful, or that disrupts civil, constructive discourse"*,

a three-level valence rating toward the film (negative, mixed, positive), a three-level stance toward seeing the film (warn, neutral, recommend), an emotionality rating (low, medium, high), and a sarcasm indicator. The GPT-4o-mini toxicity rating serves as a model-independent cross-validation of the Perspective scores. The categorical ratings enable post-hoc analyses examining whether any observed effects of toxicity dose are moderated by the specific type of toxicity that participants were exposed to.

Corpus Characteristics

The final corpus contains 461 organic posts spanning a Perspective toxicity range of [0.019, 0.945] (mean = 0.459, SD = 0.268). Applying the classification threshold of 0.6 yields 300 civil posts (65.1%) and 161 toxic posts (34.9%), providing sampling pools of sufficient depth for the dose-response design. Figure 5 shows the Perspective score distribution across the full corpus. Table 4 shows illustrative posts from five toxicity bands.

Convergent validity. The Perspective API score and the GPT-4o-mini rating are correlated at $r = .75$ (95% bootstrap CI [.71, .79]; $N = 388$ posts), indicating that the two automated measures converge on the same underlying construct despite relying on independent model families and training objectives. Figure 6 displays the joint distribution.

Potential confounds. A concern with any constructed stimulus corpus is that the intended manipulation — toxicity — may co-vary with unintended text features that could independently affect participant responses. The two most plausible confounds are post length and readability. We examined both across the ten toxicity bins. Post length (word count) showed no systematic relationship with toxicity (Spearman $\rho = .04$), confirming that longer posts are not disproportionately concentrated in any part of the toxicity range. Readability (Flesch-Kincaid grade level) was likewise approximately uniform across bins (see Table 3). Beyond these surface features, all posts share a common topic (Avatar: Fire and Ash), were generated from the same prompt structure, and cover a comparable range of aspects (plot, cast, visuals, comparisons to prior instalments). Systematic variation in toxicity level was therefore not accompanied by systematic variation in topic, framing, or linguistic complexity.

Human Validation of Toxicity Scores

Design

To establish the criterion validity of both automated toxicity measures – the Perspective API score and the GPT-4o-mini rating – we collected human toxicity ratings

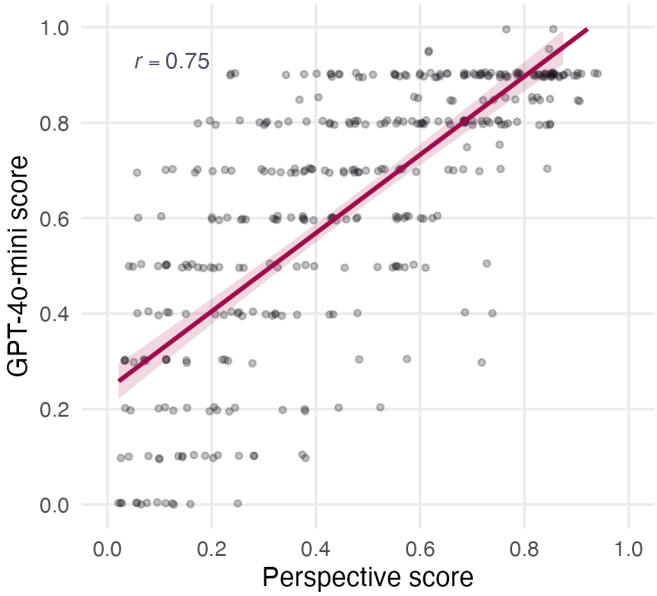

Figure 6. Convergent validity of the two automated toxicity measures. Each point is one stimulus post. The line is an OLS fit. Points are jittered slightly to reduce overplotting.

Table 3. Post length and readability by toxicity bin. Word count and Flesch-Kincaid grade level (FK-GL) are averaged within each bin. Absence of a monotone trend indicates that toxicity does not co-vary with these surface features.

Bin	N	Words M	Words SD	FK-GL M	FK-GL SD
[0–0.1)	50	31.0	15.9	15.0	2.2
[0.1–0.2)	50	32.4	15.8	15.0	2.3
[0.2–0.3)	50	28.8	12.7	14.9	2.0
[0.3–0.4)	50	32.6	15.0	15.1	2.5
[0.4–0.5)	50	35.9	16.4	15.5	1.9
[0.5–0.6)	50	29.1	12.6	15.4	2.1
[0.6–0.7)	50	31.6	18.0	16.4	2.9
[0.7–0.8)	50	37.3	17.6	15.7	1.8
[0.8–0.9)	50	38.0	18.3	15.0	1.9
[0.9–1)	11	17.9	7.4	14.6	2.3

from a separate sample of Prolific participants. We randomly drew 100 posts from the stimulus corpus, stratified across the ten Perspective score bins (10 posts per bin) to ensure coverage across the full toxicity range. Pure random sampling would undersample the sparse upper tail and provide no validation signal where classifier accuracy is hardest to establish.

The study was implemented in oTree (Chen, Schonger, and Wickens 2016) and delivered to participants via Prolific. Eligibility was restricted to participants with a 99% or higher task-approval rate and self-reported fluency in English. Each participant rated all 100 posts. Posts were presented in randomized order across participants, 10 per page, with the toxicity definition displayed as a fixed header on every page. Participants moved a continuous slider (0–100) to indicate their rating for each post before advancing. The slider had no default position, requiring an active response for every item.

Toxicity was defined for raters as follows:

> *"language that a reasonable person would find attacking, devaluing, or harmful – or that disrupts civil, constructive discourse."*

This definition was also used for the GPT-4o-mini ratings, which ensures that the human and GPT ratings share a common construct definition and that any correlation

Table 4. Illustrative stimulus posts sampled from five toxicity bands. Perspective score is the Google Perspective API toxicity score; GPT score is the model-independent GPT-4o-mini rating.

Perspective	GPT	Post text
0.041	0.50	It's honestly sad that this movie pulled $1.485B. Three reasons Disney keeps milking this franchise no matter how repetitive: Quaritch recycled yet again, cheaper sequels on the way, and fans getting sold the same story. Open your eyes, people, it's all about the cash.
0.202	0.60	Gotta admit, the IMAX visuals in this one are jaw-dropping, but they don't rescue the overstuffed mess. If you think it improves on rewatch, you're fooling yourself. Varang's story in the next flick will probably be another tired trope. Snap out of it, folks.
0.480		Hey, anyone bashing Avatar 3, watch out. Trash this masterpiece, and we're gonna bury your lame takes. Bring it on!
0.667	0.85	Gimme a break, here's why the new Avatar sucks: Spider's insufferable, action sequences are garbage, and even the visuals can't salvage this mess. Let's tank this junk with bad reviews till it's gone!
0.863	0.90	Listen up, Avatar stans, you're clueless for loving this film! The fire clan battles were utter trash, complete rubbish! Go to hell for backing this mess, you useless fools! Boycott this garbage right now!

between them reflects shared construct coverage rather than shared operationalization with the Perspective API.

Three attention-check posts were embedded at fixed positions within the randomized sequence. Two were criterion items with near-universal expected responses: one unambiguously civil post (expected rating near 0) and one unambiguously toxic post (expected rating near 100). Participants who rated the civil check above 40 or the toxic check below 60 were excluded prior to analysis. A third item at an intermediate toxicity level served as a diagnostic check on rating consistency but was not used as an exclusion criterion.

Analysis

The primary validation statistics are the Pearson correlations between the mean human rating per post (averaged across retained raters) and, respectively, the Perspective API score and the GPT-4o-mini score. A higher correlation indicates better alignment between the automated measure and the human construct. We report both correlations alongside their 95% bootstrap confidence intervals (1,000 resamples at the post level).

Inter-rater reliability is reported as Krippendorff's α for ordinal data, computed across all retained raters and all 100 posts. Values above 0.80 indicate strong agreement; values between 0.67 and 0.80 are considered acceptable for exploratory purposes.

We also conduct the discrimination check recommended by Barberá et al. (2020): posts are assigned to deciles of the Perspective score distribution and the mean human rating within each decile is plotted. A well-calibrated automated measure should produce a monotone increasing relationship. Deviations from monotonicity identify score ranges where the classifier diverges from human judgment.

Results

Validation data will be collected prior to submission and results reported here.